

# Introduction to Deep Learning

AAG Hawaii meeting

April 18, 2024

Benoit Parmentier

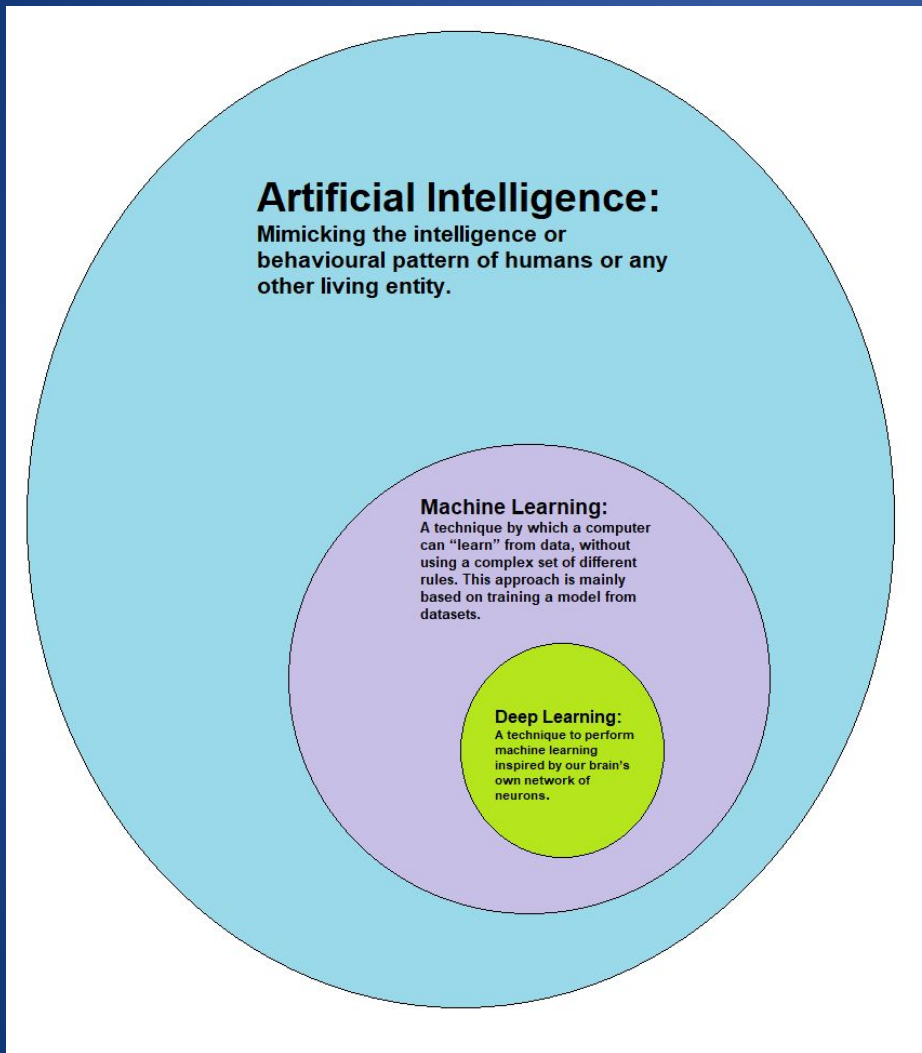
University of Mary Washington, VA

[benoit.parmentier@gmail.com](mailto:benoit.parmentier@gmail.com)

<https://www.benoitparmentier.org/>

<https://github.com/bparment1>

# Deep Learning



- Artificial Intelligence
- Machine Learning
- Deep Learning

- Image from wikipedia

# Deep Learning Key moments

- 1967: multi layers perceptron model published
- 1986: Rumelhart Hinton and Williams popularized backpropagation in a paper.
- 1989 backpropagation to learn kernels in CNN for handwritten zip code recognition
- 1998: Lenet5 architecture inspiration for many later CNNs (7 layers).
- 2006: Hinton et al. introduce methods to train NN with many hidden layers (deep networks).
- **2012:** Alexnet wins the Imagenet competition, data augmentation is introduced
- **2013:** R- CNN published for object detection
- **2015: Unet Architecture for semantic segmentation**

# TYPES OF DATA AND INPUTS IN DEEP LEARNING NETWORK

- **structured (tabular):**  
in the form of a table/data frame with column corresponding to the features/covariates
- **text:**  
string of characters with words and sentence. Although, it is sometime called unstructured, word association and sequences contain pattern and structure.
- **images:**  
matrix of width and height, also called unstructured data but contains spatial configuration that can be leveraged in deep learning networks.
- **networks:**  
matrix of width and height, also called unstructured data but contains spatial configuration that can be leveraged in deep learning networks.
- **sound:**  
matrix of width and height, also called unstructured data but contains spatial configuration that can be leveraged in deep learning networks.
- **others:**  
biochemistry, biology etc.

# Deep Learning Key Players

- **Yann Lecun:**  
French computer scientist, contributed to CNN with learned kernels. Famous for the Lenet network for digit classification.
- **Geoffrey Hinton:**  
based at the University of Toronto Computer Science department, early pioneer. Contributed to many incremental improvements that led to the Deep Learning revolution. One of the author of Alexenet.
- **Yoshua Bengio:**  
based at Montreal Institute of Machine Learning, early contributor via Theano library.
- **Ian Goodfellow:**  
Inventor of GAN (Generative Adversarial Network) now works at Google.


## Key Terminology in relation to Earth Observation/RS

- **Image classification:** Assigning a label to a whole image (not at a pixel level).
- **Object localization:** Assigning an Image Label and bounding box for an object (e.g. cat). This assumes that there is only one object of one class present in the image.
- **Object Detection:** assigning labels and bounding boxes to many objects in images.
- **Semantic Segmentation:** assigning a class to every pixel in an image.
- **Instance Segmentation:** assigning a class and item number to object (e.g. car1, car2, bike1, bike2)




# Key Terminology in relation to Earth Observation/RS

- **height:** number of rows
- **width:** number of columns
- **channels:** RS bands
- **samples:** number of images



<https://livebook.manning.com/concept/deep-learning/shape>

# DNN FOR CLASSIFICATION/REGRESSION



- Input layer has a number of nodes equal to features
- Hidden layers are the layers between input and output.
- Output layer has the number of nodes equal to the number of classes in classification or dimensions in regression.




# DNN FOR CLASSIFICATION ON IMAGES

```
model_dnn = keras.Sequential()  
model_dnn.add(layers.Flatten(input_shape=x_train.shape[1:]))  
model_dnn.add(layers.Dense(512,activation='relu'))  
model_dnn.add(layers.Dense(10,activation='softmax'))  
  
model_dnn.compile(loss='categorical_crossentropy',  
                  optimizer='adam',  
                  metrics=['accuracy'])  
  
model_dnn._name = 'model_dnn' #set keras model name
```

- For multi-class classification, the final activation layer is softmax. Scores sum to 1 and can be interpreted as probabilities.
- Note that the categorical cross entropy is used as a loss. This is the case for multi-class classification.

# CONVOLUTIONAL LAYER



[https://miro.medium.com/max/2880/0\\*QS1ArBEUJjySXhE.png](https://miro.medium.com/max/2880/0*QS1ArBEUJjySXhE.png)

- A layer that uses kernel and generates feature maps.
- The weights in the kernel are learned by the network.
- It reduces greatly the number of parameters compared to a Dense layer and leverage spatial information/structure.
- <https://poloclub.github.io/cnn-explainer/>

# CNN EXPLAINER



## What is a Convolutional Neural Network?

In machine learning, a classifier assigns a class label to a data point. For example, an *image classifier* produces a class label (e.g. bird, plane) for what objects exist within an image. A *convolutional neural network*, or CNN for short, is a type of classifier, which excels at solving


- <https://poloclub.github.io/cnn-explainer/>

# CONVOLUTIONAL LAYER PARAMETERS

- kernel size: Height x Width typically 3x3 or 5x5.
- stride: steps between sliding of the kernels. Typically 1. If greater than one the size of the image may decrease more.
- padding: padding added when modeler wants to keep the input size of the image.
- filter: in keras, it is the number of feature maps (generated by the equivalent number of kernels)



# CNN FOR CLASSIFICATION



- A typical sequence in a CNN: Convolutional 2D followed by max pooling.
- The network head is a dense neural network (fully connected).

# CNN FOR CLASSIFICATION

```
num_classes =10
model_cnn = Sequential()

## CNN part
#Block 1
model_cnn.add(Conv2D(32, (3, 3), padding='same',
                    activation='relu',
                    input_shape=x_train.shape[1:]))
model_cnn.add(Conv2D(32,(3,3),padding='same', activation='relu'))
model_cnn.add(MaxPooling2D(pool_size=(2,2)))


#Block 2
model_cnn.add(Conv2D(64,(3,3),padding='same',activation='relu'))
model_cnn.add(Conv2D(64,(3,3),padding='same',activation='relu'))
model_cnn.add(MaxPooling2D(pool_size=(2,2)))

#Block 3:
model_cnn.add(Conv2D(128,(3,3),padding='same',activation='relu'))
model_cnn.add(Conv2D(128,(3,3),padding='same',activation='relu'))
model_cnn.add(MaxPooling2D(pool_size=(2,2))) #becomes 4x4 image

## DNN part
model_cnn.add(Flatten())
model_cnn.add(Dropout(0.2))
model_cnn.add(Dense(1024,activation='relu',kernel_constraint=maxnorm(3)))
model_cnn.add(Dropout(0.2))
model_cnn.add(Dense(num_classes, activation='softmax'))
```



# UNET FOR SEMANTIC SEGMENTATION



- Encoder-Decoder structure
- Output is an image of the same size with probabilities for every pixel.

# NETWORK ARCHITECTURES

Feature extractor  
(CNN + Pooling)



Classifier  
DNN

- Image classification

Feature extractor  
(CNN + Pooling)



Classifier  
DNN

Regression  
DNN

- Object localization

Feature extractor  
(CNN + Pooling)



Bottleneck




Upsampling/transpose  
convolution

- Semantic segmentation

# Deep Learning Key Software Libraries

- Theano: one of the early Deep Learning library created by MILA (University of Montreal) by team led by Yoshua Bengio.
- Tensorflow: created a google for internal research in AI, was open sourced in 2015. It is now one of the main library used in industry and in research.
- Keras: created in 2015 as a high level API sitting on top of Theano, Tensorflow and CNTK. It is now one of most popular Deep Learning libraries.
- **pytorch**: created by facebook AI, now one of the main library used in industry and research.
- **fastai**: high level libraries that sits on top of pytorch with the goal of making Deep Learning accessible to all. Developed at the University of San Francisco.

# Training a deep learning model



# Compiling model

```
#NOTE INPUT SHOULD BE THE NUMBER OF VAR
#### Test with less number of input nodes: pruning
model_dnn = Sequential()
model_dnn.add(Dense(4, input_dim=3, activation='relu'))
model_dnn.add(Dense(10, activation='relu'))
model_dnn.add(Dense(10, activation='relu'))
model_dnn.add(Dense(1)) #scalar regression, end DNN without activation function as we are predicting continuous values

model_dnn.compile(optimizer="adam",
                  loss="mse",
                  metrics=["mae"])
```

```
model_dnn.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 4)	16
dense_1 (Dense)	(None, 10)	50
dense_2 (Dense)	(None, 10)	110
dense_3 (Dense)	(None, 1)	11

```
Total params: 187
```

```
Trainable params: 187
```

```
Non-trainable params: 0
```

- Once a model is created, you need to compile the model and select the optimizer, loss function and metrics.

# Loss Function

- A loss function is an error function that is optimized in the model.
- It is function of the weights and it is used to find weight values to update.
- Gradient descent is used to find the minimum of the loss function and the optimal weights.
- Loss function types varies in function of input data: continuous (e.g. MSE) or categorical (e.g. binary cross entropy).



# Loss Function

- MSE: Mean Square Error is used as loss function for neural networks with a continuous target (regression).
- Binary cross entropy is used for binary classification or multilabel cases.
- Categorical cross-entropy is used for multi class cases (mutually exclusive classes).

$$L_{CE} = - \sum_{i=1}^n t_i \log(p_i), \text{ for } n \text{ classes,}$$

where  $t_i$  is the truth label and  $p_i$  is the Softmax probability for the  $i^{th}$  class.

<https://towardsdatascience.com/cross-entropy-loss-function-f38c4ec8643e>

# Optimizer

- It is the method used to update the weights so that we decrease the loss (error).
- It is usually a variation of the the gradient descent method to reach the minimum of the loss function.
- There are many optimizers. Some common ones are SGD, RMSprop and Adam.

# Optimizers

- **SGD**: gradient descent that is updated by batch (mini-batch gradient descent).
- **SGD with momentum**: consider previous gradient values to carry out current weight update.
- **SGD with scheduled learning rate**: pre-defined decrease in rate of learning.
- **Adagrad**: scheduled learning called Adaptive learning rates for each input features based on update frequency for each input features.
- **Rmsprop**: improvement of Adagrad with better update of learning rates.
- **Adams**: consider both scheduled learning rate and momentum trying to combine all the improvements together in one algorithm.

# The importance of Labeling

- . Deep Learning requires a large amount of labeled data training.
- . These are usually hand-labeled datasets that are expensive and time consuming to generate: <https://www.bbc.com/news/technology-56414491>
- . There are softwares to help.
- . New paradigm: weak learning.

# LABELING BY HUMAN

## AI: Ghost workers demand to be seen and heard

By Jane Wakefield  
Technology reporter

🕒 4 days ago



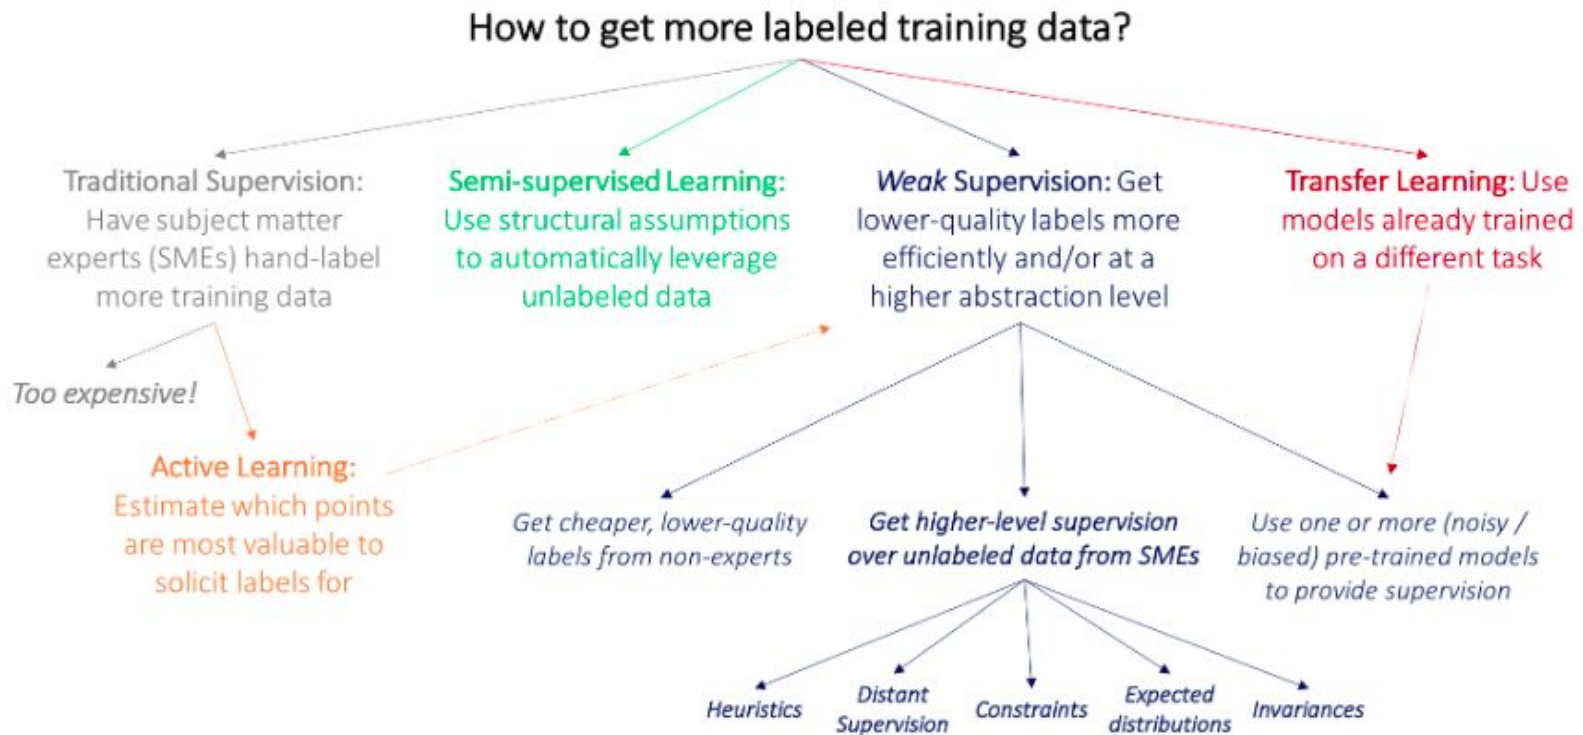
**Artificial intelligence and machine learning exist on the back of a lot of hard work from humans.**

Alongside the scientists, there are thousands of low-paid workers whose job it is to classify and label data - the lifeblood of such systems.

But increasingly there are questions about whether these so-called ghost workers are being exploited.



# The importance of Labeling





# Transfer Learning and Pre-training

- Use model trained on another related task to predict on the task at hand.
- Can leverage networks that were trained on a large training dataset.
- Often requires removing the 'head' of the model to use the 'base'.
- This method has proven very successful in many instances.

# Introduction to Deep Learning

AAG Hawaii meeting

April 18, 2024

Benoit Parmentier

University of Mary Washington, VA

[benoit.parmentier@gmail.com](mailto:benoit.parmentier@gmail.com)

<https://www.benoitparmentier.org/>

<https://github.com/bparment1>

# Introduction to Deep Learning

University of Mary and Washington

April 11, 2024

Benoit Parmentier

University of Mary Washington, VA

[benoit.parmentier@gmail.com](mailto:benoit.parmentier@gmail.com)

<https://www.benoitparmentier.org/>

<https://github.com/bparment1>

# References

- <https://ai.facebook.com/blog/billion-scale-semi-supervised-learning/>
- <https://www.azavea.com/blog/2020/03/24/labeling-satellite-imagery-for-machine-learning/>
- <http://ai.stanford.edu/blog/weak-supervision/>
- [https://www.d2l.ai/chapter\\_prelude/index.html](https://www.d2l.ai/chapter_prelude/index.html)
-